

Stock Predictions Using Machine Learning & Python

Xuan He

Alex Koynoff

Prathyusha Challa

Zelalem Kebede

Itunu Oyeyipo

Project Overview

- Use Machine Learning to predict stock prices
- ML models selected:
 - scikit-learn(sklearn) - Linear Regression
 - scikit-learn(sklearn) - Support Vector Regression (SVR)
- Picked 10 random stocks and loaded 12 months of historical prices
- Predicted 30 trading days and compared to the real prices
 - Simulations for each model ran over 200 times. Took the average of all simulations for the visualization
- Tools used:
 - Excel
 - Pandas
 - Matplotlib
 - Tableau
- Review and visualize the results
- Summary and commentary

Machine Learning Models Used

- **Sklearn - Linear Regression (LR)** (from `sklearn.linear_model` import `LinearRegression`)

Overview

- Searching for a relationship among variables
 - Dependent variables (y) or responses
 - Independent variables (x) or predictors

Pros

- One of the most widely used as it is easier to interpret the results
- Works on any size of dataset
- Gives information about relevance features

Cons

- Conversely, it assumes there is a straight-line relationship between dependent and independent variables which is sometimes incorrect
- It can be very sensitive to outliers or anomalies
- The algorithm assumes data is normally distributed in real when they are not.
- It is not recommended for most practical applications because it oversimplifies real world problems.



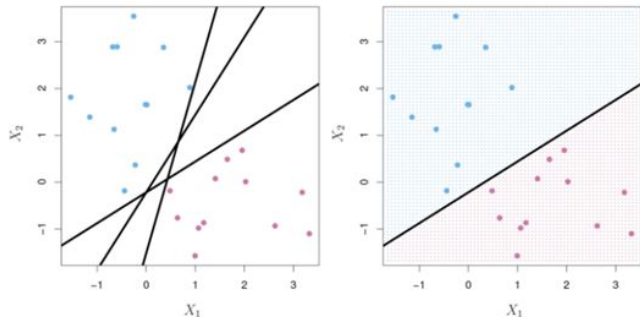
Machine Learning Models Used

- **Sklearn - Support Vector Machine (from sklearn.svm import SVR)**

Overview

- A linear model for classification and regression problems
- SVM algorithm for pattern recognition
- The algorithm creates a line or hyperplane which separates the data into classes
- The basic idea of SVM:
 - Just like 1-layer or multi-layer neural nets

Separating Hyperplanes



Machine Learning Models Used

- **Sklearn - Support Vector Machine (from sklearn.svm import SVR) cont:**

Pros

- Works well on smaller cleaner datasets
- It can be more efficient because it uses a subset of training points
- Helps analyse data and identify trends.
- Compute the multiplication of independent variables
- Require less training when dealing with smaller datasets
- No distribution requirements

Cons

- Isn't suited to larger datasets as the training time with SVMs can be high
- Less effective on noisier datasets with overlapping classes
- Can be inefficient to train, memory intensive and annoying to run and tune

```
quandl.ApiConfig.api_key = 'qKYYca8_q3vn5ws7Fpwz'
stock_list = ["TSLA", "MSFT", "FB", "WMT", "DG", "JPM", "DOV", "XOM", "KO", "MMM"]
```

```
for x in stock_list:

    temp_stock = "WIKI/" + x
    quandl_stock_list.append(temp_stock)
    df = quandl.get(temp_stock, trim_start = "2017-03-01", trim_end = "2018-03-27")
    df = df[['Adj. Close']]
    i = 0

    for index, row in df.iterrows():

        temp_close_price = str(round(row['Adj. Close'],2))
        temp_date = str(df.index[i])
        temp_stock = x
        i = i + 1

        temp_date_list.append(temp_date)
        temp_close_price_list.append(temp_close_price)
        temp_stock_list.append(temp_stock)

    stock_close_df = pd.DataFrame({'Stock Name': temp_stock_list})
    stock_close_df['Date'] = temp_date_list
    stock_close_df['Close Price'] = temp_close_price_list

    # Start predict: forecast_out will be the number of days to predict
    forecast_out = 30
    df['Prediction'] = df[['Adj. Close']].shift(-forecast_out)

    # Define X data set: Convert the dataframe to a numpy array
    # Remove the last forecast_out row
    X = np.array(df.drop(['Prediction'],1))
    X = X[:-forecast_out]

    # Define y data set:
    # Get all the y values except the last forecast_out row
    y = np.array(df['Prediction'])
    y = y[:-forecast_out]

    # Split the data into 80% training and 20% testing
    x_train, x_test, y_train, y_test = train_test_split(X, y, test_size=0.2)

    # Create and train the Support Vector Machine (Regressor)
    svr_rbf = SVR(kernel='rbf', C=1e3, gamma=0.2)
    svr_rbf.fit(x_train, y_train)

    # Testing Model: Score returns the coefficient of determination R^2 of the prediction.
    # The best possible score is 1.0
    svm_confidence = svr_rbf.score(x_test, y_test)

    # Create and train the Linear Regression Model
    lr = LinearRegression()
    # Train the model
    lr.fit(x_train, y_train)

    lr_confidence = lr.score(x_test, y_test)
    x_forecast = np.array(df.drop(['Prediction'],1))[:-forecast_out:]

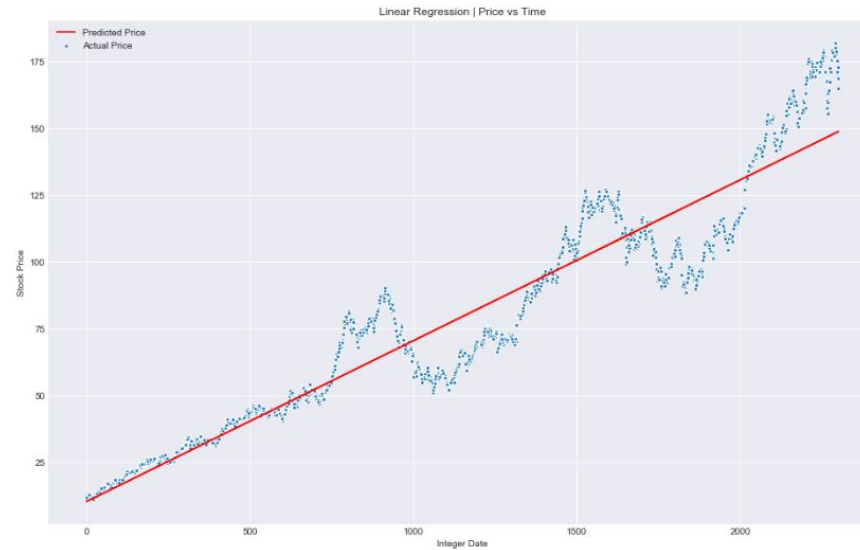

    # Prediction
    lr_prediction = lr.predict(x_forecast)
    svm_prediction = svr_rbf.predict(x_forecast)




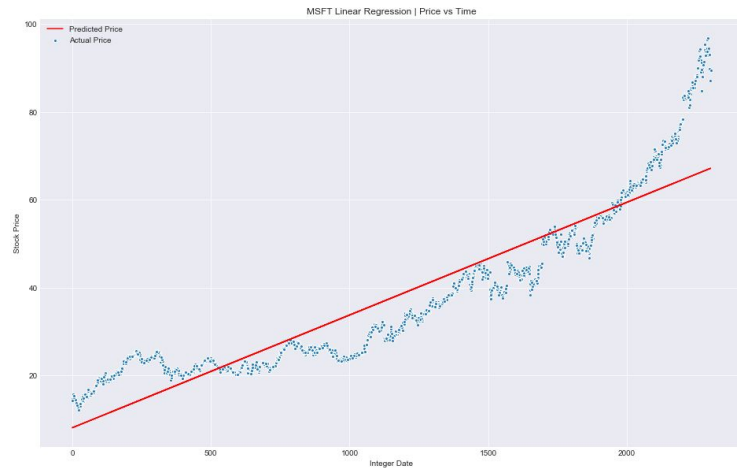
    lr_prediction_list.append(lr_prediction)
    svm_prediction_list.append(svm_prediction)

    stock_close_df["LR_Prediction"] = lr_prediction_list
    stock_close_df["SVM_Prediction"] = svm_prediction_list
```

Code

- Selected 10 random stocks
- Ran 12 months of closing prices from Quandl API
- Looped through the list of stocks
 - Selected 30 days of prediction
 - Each model included within the loop
 - Linear
 - SVR (Support Vector Regression)
 - Information put into lists and exported to csv for further manipulation and Tableau visuals.





Visualization

CLICK ME!!!

Limitations and Future Considerations

- Limitations:
 - Only one variable used (historical closing prices)
 - Probably neither Linear Regression or SVR is a good enough model to predict stock prices
- Future Considerations:
 - Add other factors that impact stock prices to be built into the model
 - News
 - Industry specific trends
 - Company specific information
 - Fundamental analysis
 - Company financials, ratios, etc
 - Technical analysis
 - Moving averages, price patterns, charting techniques, etc
 - Build and test with other models and compare results
 - Deep Learning - Long Short Term Memory Network (LSTM)

Summary

- **Do not depend on our model to trade stocks! You will lose money!**
- **Do not totally trust any prediction online! Make your own decision!**
- Add more variables for consideration in models to make them more accurate
- Use different models and compare results
- Use at your own risk!

Q & A

(or we can take a look at what ML can do and see few top trending stocks suggested by ML)



*Thank
you!*



What Machine Learning Can and Cannot Do ?

- Learning a function that maps well-defined inputs to well-defined outputs
- Large (digital) data sets exist or can be created containing input-output pairs
- No long chains of logic or reasoning that depend on diverse background knowledge or common sense
- The task provides clear feedback with clearly definable goals and metrics
- No need for detailed explanation of how the decision was made
- A tolerance for error and no need for provably correct or optimal solutions
- The phenomenon or function being learned should not change rapidly over time

THE WALL STREET JOURNAL.

Five of the Best Stocks to buy based on **US News**



- Advanced Micro Devices (AMD)
- Chipotle Mexican Grill (CMG)
- Nike (NKE)
- Canopy Growth Corp. (CGC)
- Microsoft Corp. (MSFT)

***** Disclaimer : If you happen to make money buying the suggested stocks, Venmo here ⇒***

